

ChromaDB and Retrieval Augmented Generation

ChromaDB is a vector database used to store and retrieve information using embeddings.

Embeddings represent text as numerical vectors.

Similar documents can be found using vector similarity.

Retrieval Augmented Generation, or RAG, combines information retrieval with a large language model.

A RAG system first retrieves relevant documents.

The retrieved information is then provided to the language model as context.

Ollama allows large language models to run locally.

Llama models can generate answers based on retrieved information.